

student_intervention

April 11, 2016

1 Project 2: Supervised Learning

1.0.1 Building a Student Intervention System

1.1 Classification vs Regression

Your goal is to identify students who might need early intervention - which type of supervised machine learning problem is this, classification or regression? Why?

This is classification problem. We need to classify students into two groups: needed or not needed early intervention (two-class classification), hence classification task.

1.2 Exploring the Data

Let's go ahead and read in the student dataset first.

To execute a code cell, click inside it and press **Shift+Enter**.

```
In [1]: # Import libraries
import numpy as np
import pandas as pd
from sklearn.cross_validation import StratifiedShuffleSplit
from sklearn.cross_validation import KFold
from sklearn.utils import shuffle
import time
from sklearn.metrics import f1_score, make_scorer
from sklearn import tree
from sklearn import svm
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.cross_validation import train_test_split
from sklearn.grid_search import GridSearchCV
import random

In [2]: #Read student data
student_data = pd.read_csv("student-data.csv")
print "Student data read successfully!"
# Note: The last column 'passed' is the target/label, all other are feature columns
```

Student data read successfully!

Now, can you find out the following facts about the dataset? - Total number of students - Number of students who passed - Number of students who failed - Graduation rate of the class (%) - Number of features
Use the code block below to compute these values. Instructions/steps are marked using **TODOs**.

```
In [3]: # TODO: Compute desired values - replace each '?' with an appropriate expression/function call
n_students = student_data.shape[0]
n_features = student_data.shape[1] - 1
```

```

n_passed = pd.value_counts(student_data.passed == "yes")[True]
n_failed = pd.value_counts(student_data.passed == "yes")[False]
grad_rate = float(n_passed)/(n_passed + n_failed)*100
print "Total number of students: {}".format(n_students)
print "Number of students who passed: {}".format(n_passed)
print "Number of students who failed: {}".format(n_failed)
print "Number of features: {}".format(n_features)
print "Graduation rate of the class: {:.2f}%".format(grad_rate)

```

Total number of students: 395
 Number of students who passed: 265
 Number of students who failed: 130
 Number of features: 30
 Graduation rate of the class: 67.09%

1.3 Preparing the Data

In this section, we will prepare the data for modeling, training and testing.

1.3.1 Identify feature and target columns

It is often the case that the data you obtain contains non-numeric features. This can be a problem, as most machine learning algorithms expect numeric data to perform computations with.

Let's first separate our data into feature and target columns, and see if any features are non-numeric.

Note: For this dataset, the last column ('passed') is the target or label we are trying to predict.

```

In [4]: # Extract feature (X) and target (y) columns
feature_cols = list(student_data.columns[:-1]) # all columns but last are features
target_col = student_data.columns[-1] # last column is the target/label
print "Feature column(s):-\n{}".format(feature_cols)
print "Target column: {}".format(target_col)

global X_all
global y_all

X_all = student_data[feature_cols] # feature values for all students
y_all = pd.DataFrame(student_data[target_col]) # corresponding targets/labels
print "\nFeature values:-"
print X_all.head() # print the first 5 rows

```

Feature column(s):-

['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu', 'Mjob', 'Fjob', 'reason', 'guardian']

Target column: passed

Feature values:-

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	\
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	
1	GP	F	17	U	GT3	T	1	1	at_home	other	
2	GP	F	15	U	LE3	T	1	1	at_home	other	
3	GP	F	15	U	GT3	T	4	2	health	services	
4	GP	F	16	U	GT3	T	3	3	other	other	
...											
	higher	internet	romantic	famrel	freetime	goout	Dalc	Walc	health	\	
0	...	yes	no	no	4	3	4	1	1	3	
1	...	yes	yes	no	5	3	3	1	1	3	

2	...	yes	yes	no	4	3	2	2	3	3
3	...	yes	yes	yes	3	2	2	1	1	5
4	...	yes	no	no	4	3	2	1	2	5

absences	
0	6
1	4
2	10
3	2
4	4

[5 rows x 30 columns]

1.3.2 Preprocess feature columns

As you can see, there are several non-numeric columns that need to be converted! Many of them are simply yes/no, e.g. `internet`. These can be reasonably converted into 1/0 (binary) values.

Other columns, like `Mjob` and `Fjob`, have more than two values, and are known as categorical variables. The recommended way to handle such a column is to create as many columns as possible values (e.g. `Fjob_teacher`, `Fjob_other`, `Fjob_services`, etc.), and assign a 1 to one of them and 0 to all others.

These generated columns are sometimes called dummy variables, and we will use the `pandas.get_dummies()` function to perform this transformation.

```
In [5]: # Preprocess feature columns
def preprocess_features(X):
    outX = pd.DataFrame(index=X.index) # output dataframe, initially empty
    #print outX.columns[:]
    # Check each column
    for col, col_data in X.iteritems():
        # If data type is non-numeric, try to replace all yes/no values with 1/0
        if col_data.dtype == object:
            col_data = col_data.replace(['yes', 'no'], [1, 0])
            # Note: This should change the data type for yes/no columns to int

        # If still non-numeric, convert to one or more dummy variables
        if col_data.dtype == object:
            col_data = pd.get_dummies(col_data, prefix=col) # e.g. 'school' => 'school_GP', 'school_LE3'

    outX = outX.join(col_data) # collect column(s) in output dataframe

    return outX
global X_all
global y_all
X_all = preprocess_features(X_all)
y_all = preprocess_features(y_all)
print "Processed feature columns ({}):-\n{}".format(len(X_all.columns), list(X_all.columns))
```

Processed feature columns (48):-

['school_GP', 'school_MS', 'sex_F', 'sex_M', 'age', 'address_R', 'address_U', 'famsize_GT3', 'famsize_LE3']

1.3.3 Split data into training and test sets

So far, we have converted all categorical features into numeric values. In this next step, we split the data (both features and corresponding labels) into training and test sets.

```
In [6]: # First, decide how many training vs test samples you want
num_all = student_data.shape[0] # same as len(student_data)
num_train = 300 # about 75% of the data
num_test = num_all - num_train

# TODO: Then, select features (X) and corresponding labels (y) for the training and test sets
# Note: Shuffle the data or randomly select samples to avoid any bias due to ordering in the data

X_train, X_test, y_train, y_test = train_test_split(X_all, y_all, test_size=num_test, train_size=num_train)
y_train = y_train.as_matrix().reshape((y_train.shape[0]))
y_test = y_test.as_matrix().reshape((y_test.shape[0]))

# X_train, y_train, X_test, y_test = next_batch(rs, train_size = num_train, test_size = num_test)
print "Training set: {} samples".format(X_train.shape[0])
print "Test set: {} samples".format(X_test.shape[0])

# Note: If you need a validation set, extract it from within training data
```

Training set: 300 samples

Test set: 95 samples

1.4 Training and Evaluating Models

Choose 3 supervised learning models that are available in scikit-learn, and appropriate for this problem. For each model:

- What is the theoretical $O(n)$ time & space complexity in terms of input size?
- What are the general applications of this model? What are its strengths and weaknesses?
- Given what you know about the data so far, why did you choose this model to apply?
- Fit this model to the training data, try to predict labels (for both training and test sets), and measure the F1 score. Repeat this process with different training set sizes (100, 200, 300), keeping test set constant.

Produce a table showing training time, prediction time, F1 score on training set and F1 score on test set, for each training set size.

Note: You need to produce 3 such tables - one for each model.

```
In [7]: # Train a model

def train_classifier(clf, X_train, y_train):
    print "Training {}...".format(clf.__class__.__name__)
    start = time.time()
    clf.fit(X_train, y_train)
    end = time.time()
    print "Done!\nTraining time (secs): {:.3f}".format(end - start)
    return "{:.5f}".format(end - start)

# TODO: Choose a model, import it and instantiate an object

def create_classifier(type_of = "Tree", weights = None):
    if type_of == "Tree":
        return tree.DecisionTreeClassifier(class_weight=weights) #used parameter class_weight b
    elif type_of == "SVM":
        return svm.SVC(class_weight=weights) #used parameter class_weight because dataset not b
```

```

elif type_of == "GrBoost":
    return GradientBoostingClassifier(n_estimators=100, learning_rate=.1, max_depth=3)
else:
    raise ValueError("Classifier not found", type_of)

clf = create_classifier("Tree")
train_classifier(clf, X_train, y_train)
print clf # you can inspect the learned model by printing it
#print

Training DecisionTreeClassifier...
Done!
Training time (secs): 0.003
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
    max_features=None, max_leaf_nodes=None, min_samples_leaf=1,
    min_samples_split=2, min_weight_fraction_leaf=0.0,
    presort=False, random_state=None, splitter='best')

In [8]: # Predict on training set and compute F1 score
from sklearn.metrics import f1_score

def predict_labels(clf, features, target):
    print "Predicting labels using {}".format(clf.__class__.__name__)
    start = time.time()
    y_pred = clf.predict(features)
    end = time.time()
    print "Done!\nPrediction time (secs): {:.5f}".format(end - start)
    return f1_score(target, y_pred), "{:.5f}".format(end - start)

train_f1_score, time_ = predict_labels(clf, X_train, y_train)
print "F1 score for training set: {:.3f} in {} sec".format(train_f1_score, time_)

Predicting labels using DecisionTreeClassifier...
Done!
Prediction time (secs): 0.00042
F1 score for training set: 1.000 in 0.00042 sec

In [9]: # Predict on test data
print "F1 score for test set: {}".format(predict_labels(clf, X_test, y_test))

Predicting labels using DecisionTreeClassifier...
Done!
Prediction time (secs): 0.00069
F1 score for test set: (0.72307692307692306, '0.00069')

```

1.4.1 Decision Trees

Complexity

Decision Trees in general need:

$$O(n_{samples} \times n_{features} \times \log n_{samples}).$$

to construct balanced binary tree and query time:

$$O(\log n_{samples}).$$

For the project's problem complexity is:

```
In [10]: samples_ = 100
        print "Complexity for construction O({:d}) for {:d} samples".format(int(samples_*48*np.log2(samples_)), samples_)
        print "Complexity for query O({:d}) for {:d} samples".format(int(np.log2(samples_)), samples_)
```

Complexity for construction O(31890) for 100 samples

Complexity for query O(6) for 100 samples

General applications

Decision-trees best suited to problems with following characteristics:

- Samples are represented by attribute-value pairs. Each feature takes on a small number of disjoint possible values (e.g., man, woman).
- Best for two possible output values.
- Problems with disjunctive descriptions.
- The training data may contain errors. Decision-tree learning methods are robust for errors.
- Training data may contain missing attribute data.

Strengths and weaknesses

Strengths:

- Simple to understand and to interpret, especially using graphic representation.
- White box - learned model will be explained using boolean logic.
- Data preparation not very hard. No needs data normalization, dummy and empty value filtering.
- Computation cost is relatively low - logarithmic for tree training and prediction.
- Possible validate model using statistical tests.

Weaknesses:

- Decision tree learners create biased trees if some classes dominate. Need balanced training data for every class or need equal number of samples for each class.
- Usually decision-tree learners generate overfitted models and for preventing that depth max and minimum number of samples at a leaf node are necessary.
- Small variations of data might result of completely different trees being generated.
- Learning optimal decision-tree is NP-problem. In practice heuristic methods usually used and that are not guarantee globally optimal tree (to avoid this multiple decision trees will be used with randomly sampled features and samples).

Dataset analysis

```
In [11]: def check_of_empty_values(data_set):
        #create empty Data frame for broken data (NaN or empty)
        return data_set.isnull().values.any()

        def get_labels_balance(data_set):
            #countn labels in training or testing set and return label weights
            unique, counts = np.unique(data_set, return_counts=True)
            label_weights = {}
```

```

for i in range(0, unique.shape[0]):
    label_weights[unique[i]] = float(counts[i])/data_set.shape[0]
return label_weights

print "Is Student feature dataset contains NaN or empty values ?", check_of_empty_values(X_all)
print "Is Student label dataset contains NaN or empty values ?", check_of_empty_values(y_all)
print "Weights of data balance:"
print "Training data:", get_labels_balance(y_train)
print "Testing data:", get_labels_balance(y_test)

```

```

Is Student feature dataset contains NaN or empty values ? False
Is Student label dataset contains NaN or empty values ? False
Weights of data balance:
Training data: {0: 0.3233333333333333, 1: 0.6766666666666666}
Testing data: {0: 0.3473684210526316, 1: 0.6526315789473685}

```

Justification

Decision-Trees will be used for classification and regression problems with single and multi-variable output. Since DT is white box it is possible to explain prediction results and it may be useful for stuff.

Looking on the students data the DT also may be used for such problem because:

- most of attributes values are two-pared;
- predicting value belongs for the two classes;
- dataset specially was not prepared (excluding data type conversion to numbers and using dummies variables).

Although data set is not balanced this may mitigated by using label_weights parameter for classifier.

Training and prediction

```

In [12]: # Train and predict using different training set sizes
def train_predict(clf, X_train, y_train, X_test, y_test):
    #print "-----"
    #print "Training set size: {}".format(len(X_train))
    tr_weights = get_labels_balance(y_train)
    time_str = train_classifier(clf, X_train, y_train)
    tr_pred = predict_labels(clf, X_train, y_train)
    #print "F1 score for training set: {}".format(tr_pred)
    ts_pred = predict_labels(clf, X_test, y_test)
    #print "F1 score for test set: {}".format(ts_pred)
    return [{"Size":len(X_train), "Time":time_str, "F1_training":tr_pred, "F1_testing":ts_pred}

#Create Pandas table
def append_value(frame, dict_values, i):
    if type(frame) is pd.DataFrame:
        new_frame = pd.DataFrame(dict_values, index=[i])
        return pd.concat([frame, new_frame])
    else:
        return pd.DataFrame(dict_values, index=[i])

```

```

# TODO: Run the helper function above for desired subsets of training data
# Note: Keep the test set constant
training_set_sizes = [100, 200, 300]
tree_table = None
i = 1
for set_size in training_set_sizes:
    #X_train, X_test, y_train, y_test = train_test_split(X_all, y_all, test_size=num_test, tra
    #X_train, y_train, X_test, y_test = next_batch(rs, train_size=set_size, test_size=num_test

    # Create Decision-tree classifier and use label weights
    lab_weights = get_labels_balance(y_train)
    clf = create_classifier("Tree", weights=lab_weights)
    train_time = train_classifier(clf, X_train, y_train)
    f1_training, tr_time = predict_labels(clf, X_train, y_train)
    f1_testing, ts_time = predict_labels(clf, X_test, y_test)
    row = {"Size":set_size, "F1_training":f1_training, "F1_testing":f1_testing, "Time train":t
    tree_table = append_value(tree_table, row, i)
    i += 1

# Result of Decision Tree
print "\nDecision Tree training results\n", tree_table

```

```

Training DecisionTreeClassifier...
Done!
Training time (secs): 0.003
Predicting labels using DecisionTreeClassifier...
Done!
Prediction time (secs): 0.00030
Predicting labels using DecisionTreeClassifier...
Done!
Prediction time (secs): 0.00018
Training DecisionTreeClassifier...
Done!
Training time (secs): 0.002
Predicting labels using DecisionTreeClassifier...
Done!
Prediction time (secs): 0.00025
Predicting labels using DecisionTreeClassifier...
Done!
Prediction time (secs): 0.00017
Training DecisionTreeClassifier...
Done!
Training time (secs): 0.002
Predicting labels using DecisionTreeClassifier...
Done!
Prediction time (secs): 0.00020
Predicting labels using DecisionTreeClassifier...
Done!
Prediction time (secs): 0.00015

```

Decision Tree training results

	F1_testing	F1_training	Size	Time test	Time train
1	0.714286	1	100	0.00018	0.00286
2	0.714286	1	200	0.00017	0.00235

3	0.698413	1	300	0.00015	0.00223
---	----------	---	-----	---------	---------

Summary

Training time increase with approximately same rate as training size increase. Testing time fluctuating on the same level. Model for all training sets looks like overfitted.

1.4.2 Support Vector Machines

Complexity. The core of SVM is quadratic programming problem (QP), separating support vectors from the rest of the training data. Support Vector Machines for scipy implementation needs between

$$O(n_{features} \times n_{samples}^2)$$

and

$$O(n_{features} \times n_{samples}^3)$$

For the project's problem complexity is:

```
In [13]: samples_ = 100
          print "Complexity for construction O({:.3e}) for {:d} samples".format(int(48*samples_**2), samples_)
          print "Complexity for query O({:.3e}) for {:d} samples".format(int(48*samples_**3), samples_)
```

Complexity for construction O(4.800e+05) for 100 samples

Complexity for query O(4.800e+07) for 100 samples

General applications. Support vector machines (SVMs) are a set of supervised learning methods used for classification, regression and outliers detection. Effective in high dimensional spaces. Still effective in cases where number of dimensions is greater than the number of samples. Uses a subset of training points in the decision function (called support vectors), so it is also memory efficient.

Strengths and weaknesses

Strengths:

- Effective in high dimensional spaces.
- Still effective in cases where number of dimensions is greater than the number of samples.
- Uses a subset of training points in the decision function (called support vectors), so it is also memory efficient.
- Versatile: different Kernel functions can be specified for the decision function. Common kernels are provided, but it is also possible to specify custom kernels.

Weakness:

- If the number of features is much greater than the number of samples, the method is likely to give poor performances.
- SVMs do not directly provide probability estimates, these are calculated using an expensive five-fold cross-validation (see Scores and probabilities, below).
- Support Vector Machine algorithms are not scale invariant.

Justification

- Number of features is relatively large compare to number of samples.
- SVM will be used for two-class problem.
- Prediction will use relatively small piece of memory and be quick- only support vectors are stored.
- SVM also will be used for classification problems.
- Kernel function can express domain knowledge.

Training and prediction

```
In [14]: # SVM
training_set_sizes = [100, 200, 300]
svm_table = None
i = 1
for set_size in training_set_sizes:

    #X_train, X_test, y_train, y_test = train_test_split(X_all, y_all, test_size=num_test, tra
    #X_train, y_train, X_test, y_test = next_batch(rs, train_size=set_size, test_size=num_test

    # Create Decision-tree classifier and use label weights
    lab_weights = get_labels_balance(y_train)
    clf = create_classifier("SVM", weights=lab_weights)

    train_time = train_classifier(clf, X_train, y_train)
    f1_training, tr_time = predict_labels(clf, X_train, y_train)
    f1_testing, ts_time = predict_labels(clf, X_test, y_test)
    row = {"Size":set_size, "F1_training":f1_training, "F1_testing":f1_testing, "Time train":t
    svm_table = append_value(svm_table, row, i)
    i += 1

# SVM training results
print "\nSVM training results\n", svm_table
```

Training SVC...

Done!

Training time (secs): 0.011

Predicting labels using SVC...

Done!

Prediction time (secs): 0.00929

Predicting labels using SVC...

Done!

Prediction time (secs): 0.00236

Training SVC...

Done!

Training time (secs): 0.011

Predicting labels using SVC...

Done!

Prediction time (secs): 0.00745

Predicting labels using SVC...

Done!

Prediction time (secs): 0.00266

Training SVC...

Done!

Training time (secs): 0.010

```
Predicting labels using SVC...
Done!
Prediction time (secs): 0.00817
Predicting labels using SVC...
Done!
Prediction time (secs): 0.00249
```

SVM training results

	F1_testing	F1_training	Size	Time test	Time train
1	0.789809	0.807157	100	0.00236	0.01111
2	0.789809	0.807157	200	0.00266	0.01082
3	0.789809	0.807157	300	0.00249	0.00982

Summary. The training time very rapidly increase with increasing training set. For 100 is 0.00232, but for 300 is 0.00909 or:

```
In [15]: print int(0.00909/0.00232), "times"
```

3 times

Testing time increase but with less rate as training time.

1.4.3 Gradient Boosting

Complexity. The algorithm for Boosting Trees evolved from the application of boosting methods to regression trees. The general idea is to compute a sequence of (very) simple trees, where each successive tree is built for the prediction residuals of the preceding tree. The complexity of Gradient Boosting depends on number of decision trees and their depth and features number.

General applications. Gradient Boosted Regression Trees (GBRT) is a generalization of boosting to arbitrary differentiable loss functions. GBRT is an accurate and effective off-the-shelf procedure that can be used for both regression and classification problems.

Strengths and weakness

Strengths:

- Natural handling of data of mixed type (= heterogeneous features)
- Predictive power
- Robustness to outliers in output space (via robust loss functions)
- Robustness to data scaling

Weakness:

- Scalability, due to the sequential nature of boosting it can hardly be parallelized.

Justifications

- GradientBoostingClassifier supports both binary and multi-label classification
- Prediction time relatively low
- Supports mixed type of features and not needed data normalization

Training and prediction

```
In [16]: training_set_sizes = [100, 200, 300]
         boost_table = None
         i = 1
         for set_size in training_set_sizes:
             #X_train, X_test, y_train, y_test = train_test_split(X_all, y_all, test_size=num_test, tra
             #X_train, y_train, X_test, y_test = next_batch(rs, train_size=set_size, test_size=num_test

             clf = create_classifier("GrBoost", weights=lab_weights)

             train_time = train_classifier(clf, X_train, y_train)
             f1_training, tr_time = predict_labels(clf, X_train, y_train)
             f1_testing, ts_time = predict_labels(clf, X_test, y_test)
             row = {"Size":set_size, "F1_training":f1_training, "F1_testing":f1_testing, "Time train":t
             boost_table = append_value(boost_table, row, i)
             i += 1

         # Gradient Boosting training results
         print "\nGradient Boosting training results\n", boost_table
```

```
Training GradientBoostingClassifier...
Done!
Training time (secs): 0.126
Predicting labels using GradientBoostingClassifier...
Done!
Prediction time (secs): 0.00129
Predicting labels using GradientBoostingClassifier...
Done!
Prediction time (secs): 0.00062
Training GradientBoostingClassifier...
Done!
Training time (secs): 0.138
Predicting labels using GradientBoostingClassifier...
Done!
Prediction time (secs): 0.00156
Predicting labels using GradientBoostingClassifier...
Done!
Prediction time (secs): 0.00095
Training GradientBoostingClassifier...
Done!
Training time (secs): 0.141
Predicting labels using GradientBoostingClassifier...
Done!
Prediction time (secs): 0.00132
Predicting labels using GradientBoostingClassifier...
Done!
Prediction time (secs): 0.00063
```

Gradient Boosting training results

	F1_testing	F1_training	Size	Time test	Time train
1	0.805755	0.968974	100	0.00062	0.12639
2	0.820144	0.968974	200	0.00095	0.13838
3	0.805755	0.968974	300	0.00063	0.14087

Summary Training time slowly increase as increasing training set size. Prediction time fluctuates on same level.

1.4.4 Summary of algorithms

```
In [17]: print "Decision Tree"
         print tree_table

         print "\nSVM"
         print svm_table

         print "\nGradient boosting"
         print boost_table
```

Decision Tree

	F1_testing	F1_training	Size	Time test	Time train
1	0.714286	1	100	0.00018	0.00286
2	0.714286	1	200	0.00017	0.00235
3	0.698413	1	300	0.00015	0.00223

SVM

	F1_testing	F1_training	Size	Time test	Time train
1	0.789809	0.807157	100	0.00236	0.01111
2	0.789809	0.807157	200	0.00266	0.01082
3	0.789809	0.807157	300	0.00249	0.00982

Gradient boosting

	F1_testing	F1_training	Size	Time test	Time train
1	0.805755	0.968974	100	0.00062	0.12639
2	0.820144	0.968974	200	0.00095	0.13838
3	0.805755	0.968974	300	0.00063	0.14087

1.5 Choosing the Best Model

- Based on the experiments you performed earlier, in 1-2 paragraphs explain to the board of supervisors what single model you chose as the best model. Which model is generally the most appropriate based on the available data, limited resources, cost, and performance?
- In 1-2 paragraphs explain to the board of supervisors in layman's terms how the final model chosen is supposed to work (for example if you chose a Decision Tree or Support Vector Machine, how does it make a prediction).
- Fine-tune the model. Use Gridsearch with at least one important parameter tuned and with at least 3 settings. Use the entire training set for this.
- What is the model's final F1 score?

1.5.1 Most appropriate model

The SVM was choosed as most appropriate model for the particular problem. > DecisionTrees model looks like overfitted on the given data and size of it. Although training and testing time was lowest comparing to other models. DS model was rejected - on given data it overfitting and can't give general model.

Gradient boosting for the given problem may be used, but there are no more data to construct general model. For example, of the size 300 model became overfitted. Although Gradient Boosting is white box model and this can be used for staff. Gradient boosting was taken for training greater time as other models, although testing time is lowest. Gradient boosting was rejected - more data will needed and F1 accuracy score less than for SVM.

Compare to the available data and resources the SVM is more appropriate model and may construct general model for the particular problem. SVM will be more scalable (complexity calculation determined) for future maintenance. Training and testing performance acceptable (not worst and not best) comparing to the other models.

1.5.2 SVM prediction in Layman's terms

SVM is good choice for the Student intervention problem, where high dimensionality of the data (large number of descriptive fields exists). In future this dimensionality probably will increase (especially with using e-learning platform and taking into account students activities) and SVM has strong side of classifying such data.

Main idea in a SVM is to separate data with boundaries which best separates data on different classes. For example in our problem we need to classify students into two classes: passed and not passed exam. **In the learning stage** SVM try to find best way how to separate given data for the two classes and builds mathematical model for further predictions.

The best way to understand leaning of SVM is graphical representation:




Figure 1: "SVM training"

Where x_1 and x_2 are description information about students. For example, x_1 is failures (ranged) and x_2 - family status (also ranged). Circles represents students who passed exam and squares - who are not. SVM trying to find (in this example) linear separator (line) which best separates our classes of students.

What means best separator? Here answer on the next figure:




Figure 2: "Boundaries of two classes"

Best separator is two lines (boundaries) on circles and squares (in linear case) of the two classes which are given maximal distance between lines. Circles and squares on boundaries are closest to the middle line between boundaries and are named **support vectors**.

After SVM find maximal boundaries then learning process finish and SVM is ready for using.

For some problems not possible to find best separator (called **kernel function**) using linear function. In such situations will be used non-linear separator. Example you can see in Figure 3.

After a new student come to learning we have a lot of data about him/his for prediction. In the prediction stage we need to provide preprocessed (prepared) data for the learned SVM model and as result obtain prediction of him/her exam passing.

1.5.3 Tuning SVM

As main three criteria were used for tuning SVM (SVC function):

- C - penalty parameter. It corresponds to regularize more the estimation if dataset is noisy (C value will be increased). Default=1.0.




Figure 3: “Non-linear kernel”

- Kernel function: rbf (default), linear, poly, sigmoid and custom.
- Tolerance - stopping criterion tolerance (default=1e-3).

```
In [36]: c_range = np.arange(0.8, 1.9, .05)
         tol_range = np.arange(0.0001, 0.01, 0.0005)

         # parameters used with GridSearch for SVM
         params_dic = {"C":c_range, "kernel": ['rbf', 'linear', 'poly', 'sigmoid'], "tol" : tol_range}

         # prepare CV
         cv = StratifiedShuffleSplit(y_train, n_iter = 20, random_state = 42)

         # create SVM classifier and grid_search, define scorer
         scorer = make_scorer(f1_score)
         clf = svm.SVC()
         grid_search_obj = GridSearchCV(estimator=clf, param_grid=params_dic, scoring=scorer, iid=True,
                                       cv=cv, verbose=True, n_jobs=16)

         # Fit data with GridSearch
         grid_search_obj.fit(X_train, y_train)

         #get best estimator
         best_estimator = grid_search_obj.best_estimator_

         print "Best parameters:", best_estimator
```

Fitting 20 folds for each of 1760 candidates, totalling 35200 fits

```
[Parallel(n_jobs=16)]: Done 128 tasks      | elapsed:    1.9s
[Parallel(n_jobs=16)]: Done 748 tasks      | elapsed:   17.3s
[Parallel(n_jobs=16)]: Done 1004 tasks     | elapsed:   30.8s
[Parallel(n_jobs=16)]: Done 1463 tasks     | elapsed:   42.5s
[Parallel(n_jobs=16)]: Done 2208 tasks     | elapsed:   60.0s
[Parallel(n_jobs=16)]: Done 2758 tasks     | elapsed:  01.4min
[Parallel(n_jobs=16)]: Done 4179 tasks     | elapsed:  01.9min
[Parallel(n_jobs=16)]: Done 5731 tasks     | elapsed:  02.6min
[Parallel(n_jobs=16)]: Done 7280 tasks     | elapsed:  03.3min
[Parallel(n_jobs=16)]: Done 8631 tasks     | elapsed:  04.0min
[Parallel(n_jobs=16)]: Done 10083 tasks    | elapsed:  04.8min
```

```

[Parallel(n_jobs=16)]: Done 12023 tasks      | elapsed:  5.7min
[Parallel(n_jobs=16)]: Done 14107 tasks      | elapsed:  6.8min
[Parallel(n_jobs=16)]: Done 16020 tasks      | elapsed:  7.5min
[Parallel(n_jobs=16)]: Done 18086 tasks      | elapsed:  8.5min
[Parallel(n_jobs=16)]: Done 20038 tasks      | elapsed:  9.7min
[Parallel(n_jobs=16)]: Done 22095 tasks      | elapsed: 11.0min
[Parallel(n_jobs=16)]: Done 24798 tasks      | elapsed: 12.3min
[Parallel(n_jobs=16)]: Done 27263 tasks      | elapsed: 13.7min
[Parallel(n_jobs=16)]: Done 29806 tasks      | elapsed: 15.4min
[Parallel(n_jobs=16)]: Done 33111 tasks      | elapsed: 17.1min
[Parallel(n_jobs=16)]: Done 35200 out of 35200 | elapsed: 18.2min finished

```

```

Best parameters: SVC(C=0.80000000000000004, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape=None, degree=3, gamma='auto', kernel='sigmoid',
  max_iter=-1, probability=False, random_state=None, shrinking=True,
  tol=0.0001, verbose=False)

```

1.5.4 F1 score for tuned SVM

```

In [37]: # Predict data
         y_pred = best_estimator.predict(X_test)
         result = f1_score(y_test, y_pred)
         print "F1 score: {:.3f}".format(result)

```

F1 score: 0.820

In []: